

Intro to R for Biologists

IBiS Special Topics, Fall 2021

Class 8: Oct. 18, 2021

(Bash/Quest Interlude)

Erik Andersen and Shelby Blythe

Today's class

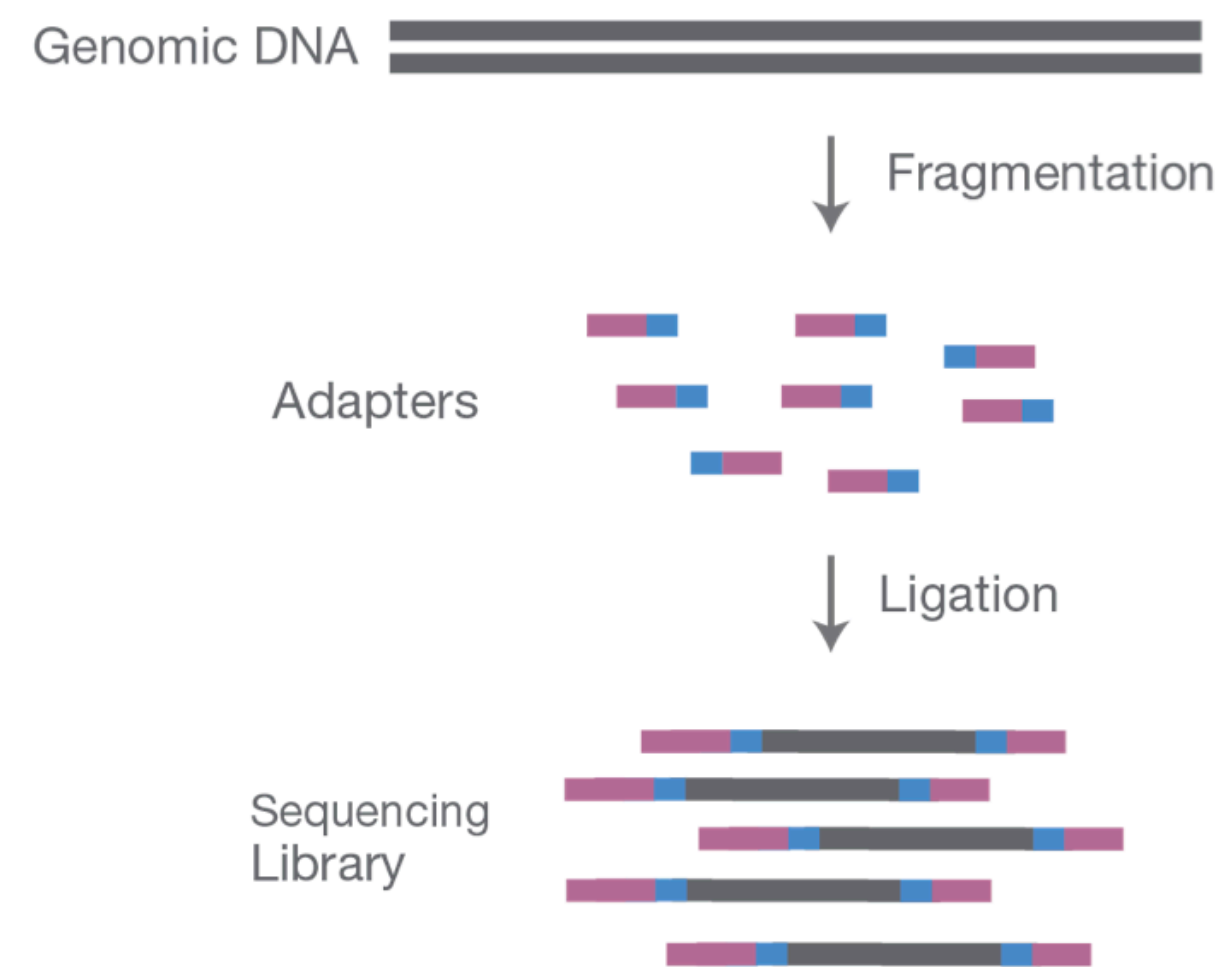
1. Illumina Sequencing and file formats
2. Fast Forward to the day you submit your paper
3. A Mapping Pipeline
4. Activity: Confirm that you can get a test script to run.

Before we begin

- Has everyone:
 1. Confirmed that they can log in to Quest with terminal and ssh and can navigate to the class allocation?
 2. Downloaded some kind of code-aware text editor (e.g., Atom or VSCode)?
 3. Successfully set up Globus (or some other appropriate file-transfer system) to send/receive files from Quest?


Illumina Sequencing

A. Library Preparation



NGS library is prepared by fragmenting a gDNA sample and ligating specialized adapters to both fragment ends.

B. Cluster Amplification



Library is loaded into a flow cell and the fragments are hybridized to the flow cell surface. Each bound fragment is amplified into a clonal cluster through bridge amplification.

The DNA fragments you sequence are capped on either end by Illumina sequencing adapters.

When you sequence the fragments, you recover X bp of sequence which may or may not be greater than the size of the original DNA fragment.

What does the data you get from the sequencer look like?

The FastQ file format:

- These are the first 4 sequences in our test data set for today's activity. They are in FastQ format.
- A single sequence is represented by 4 lines in this format.
 1. Identifier
 2. Sequence
 3. +
 4. Quality Score

```
@NS500138:372:H7K35BGXC:1:11101:25299:1039 1:N:0:TCGTCTGA+TCAAGGAC
TTCGTACTCATGGTAATGNTTATTTTCAATCTTTTA
+
AAAAEEEEEEEEEEEE#EEEEEEEEEEEEEEEE
@NS500138:372:H7K35BGXC:1:11101:7487:1040 1:N:0:TCGTCTGA+TCAAGGAC
GGGAACGGGAACGCAAGTNCACGACAAGCGCACAGG
+
AAA6EEEEEEEEEEEE#EEEE/EE6EEEEEEEE
@NS500138:372:H7K35BGXC:1:11101:2441:1040 1:N:0:TCGTCTGA+TCAAGGAC
AGGCCACATAAAGCATAANATGGCATTTCATATGCG
+
AAAAEEEEEEEEEEEE#EEEEEEEEEEEEEEEE
@NS500138:372:H7K35BGXC:1:11101:17458:1040 1:N:0:TCGTCTGA+TCAAGGAC
GTTTATATTAATAATATNCTTATAAAAAATATAAT
+
AAAAEEEEEEEEEEEE#EEEEEEEEEEEEEEEE
```

Identifier

```
@NS500138:372:H7K35BGXC:1:11101:25299:1039 1:N:0:TCGTCTGA+TCAAGGAC
```

- There is potentially useful information here:
 - The machine that did the sequencing (NS500138)
 - Run ID (372)
 - Flow Cell ID (H7K35BGXC)
 - Flow Cell Lane (1)
 - Tile (11101)
 - X-coordinate (25299)
 - Y-coordinate (1039)
 - Member of which pair? (1) *paired-end sequencing only*
 - Filtered? (N(o))
 - Control bits on? (0)
 - Barcodes (TCGTCTGA + TCAAGGAC)

Sequence and Quality Score

TTCGTACTCATGGTAATGNTTATTTTCAATCTTTTA
+
AAAAEEEEEEEEEEEEEE#EEEEEEEEEEEEEEEEEE

$$Q = -10 \log_{10} P \longrightarrow P = 10^{-\frac{Q}{10}}$$

Phred Quality Score	Probability of incorrect base call	Base call accuracy
10	1 in 10	90%
20	1 in 100	99%
30	1 in 1000	99.9%
40	1 in 10000	99.99%
50	1 in 100000	99.999%

- Each base-call is identified and brings along with it an estimate of how confident the machine is that the base-call is correct.
- The current quality score format is referred to as “Phred33”

ASCII_BASE=33 Illumina, Ion Torrent, PacBio and Sanger											
Q	P_error	ASCII	Q	P_error	ASCII	Q	P_error	ASCII	Q	P_error	ASCII
0	1.00000	33 !	11	0.07943	44 ,	22	0.00631	55 7	33	0.00050	66 B
1	0.79433	34 "	12	0.06310	45 -	23	0.00501	56 8	34	0.00040	67 C
2	0.63096	35 #	13	0.05012	46 .	24	0.00398	57 9	35	0.00032	68 D
3	0.50119	36 \$	14	0.03981	47 /	25	0.00316	58 :	36	0.00025	69 E
4	0.39811	37 %	15	0.03162	48 0	26	0.00251	59 ;	37	0.00020	70 F
5	0.31623	38 &	16	0.02512	49 1	27	0.00200	60 <	38	0.00016	71 G
6	0.25119	39 '	17	0.01995	50 2	28	0.00158	61 =	39	0.00013	72 H
7	0.19953	40 (18	0.01585	51 3	29	0.00126	62 >	40	0.00010	73 I
8	0.15849	41)	19	0.01259	52 4	30	0.00100	63 ?	41	0.00008	74 J
9	0.12589	42 *	20	0.01000	53 5	31	0.00079	64 @	42	0.00006	75 K
10	0.10000	43 +	21	0.00794	54 6	32	0.00063	65 A			

FastQ has no information about where the reads map to.

We have to process these reads and map them before they are informative.

- Typical steps in going from FastQ to “Mapped Reads”
 - Filter low-quality reads and trim adapter sequences (yields .fastq)
 - Map to reference genome (yields .sam)
 - Convert .sam to .bam (binarized .sam)
 - Sort reads (yields .bam)
 - Mark likely PCR or Optical Duplicates (yields .bam)
 - Index the .bam file (yields .bai)

The .sam/.bam format

@HD VN:1.5 S0:coordinate										
@SQ SN:ref LN:45										
r001	99	ref	7	30	8M2I4M1D3M	=	37	39	TTAGATAAAGGATACTG	*
r002	0	ref	9	30	3S6M1P1I4M	*	0	0	AAAAGATAAGGATA	*
r003	0	ref	9	30	5S6M	*	0	0	GCCTAAGCTAA	* SA:Z:ref,29,-,6H5M,17,0;
r004	0	ref	16	30	6M14N5M	*	0	0	ATAGCTTCAGC	*
r003	2064	ref	29	17	6H5M	*	0	0	TAGGC	* SA:Z:ref,9,+,5S6M,30,1;
r001	147	ref	37	30	9M	=	7	-39	CAGCGGCAT	* NM:i:1

Header
section

Alignment
section

Optional fields in the format of TAG:TYPE:VALUE

QUAL: read quality; * meaning such information is not available

SEQ: read sequence

TLEN: the number of bases covered by the reads from the same fragment. Plus/minus means the current read is the leftmost/rightmost read. E.g. compare first and last lines.

PNEXT: Position of the primary alignment of the NEXT read in the template. Set as 0 when the information is unavailable. It corresponds to POS column.

RNEXT: reference sequence name of the primary alignment of the NEXT read. For paired-end sequencing, NEXT read is the paired read, corresponding to the RNAME column.

CIGAR: summary of alignment, e.g. insertion, deletion

MAPQ: mapping quality

POS: 1-based position

RNAME: reference sequence name, e.g. chromosome/transcript id

FLAG: indicates alignment information about the read, e.g. paired, aligned, etc.

QNAME: query template name, aka. read ID

The .sam/.bam format

Example: single end mapping to *Drosophila melanogaster*

“qname” = Name of the read from fastQ	flag	rname	pos	mapq	cigar				Sequence	Phred Quality Scores
NS500138:372:H7K35BGXC:1:12111:21030:5973	16	chr2L	20634	31	36M	*	0	0	CTGGTCGGTACTAATATTTTTTCTAAGACTGTATAC	EEEE/EEEEEEEEEEEEEEEEEEEEEEEEEEEEAAAA
NS500138:372:H7K35BGXC:1:11205:21596:8754	0	chr2L	20690	1	36M	*	0	0	CGTAGGTATGTATCCGTGAGTAAGCTGAGGTGCAAA	AAAAEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEE
NS500138:372:H7K35BGXC:1:11201:23277:9037	0	chr2L	20726	1	36M	*	0	0	ACCGTATCACCATTGCTTTTCGGCTGATAACCAAT	AAAAEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEE
NS500138:372:H7K35BGXC:1:11211:25328:12582	0	chr2L	20760	1	33M	*	0	0	ATACCAAGTAAGTCAACTGTTTCCCAAATGTAT	AAAAEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEE
						rnext	pnext	tlen		

- Note, rnext, pnext, and tlen are not recorded for single-end mappings
- Are these good mappings? Two clues are the **Flag** and the **mapq** columns.

The .sam/.bam format

Example: single end mapping to *Drosophila melanogaster*

"qname" = Name of the read from fastQ	flag	rname	pos	mapq	cigar					Sequence	Phred Quality Scores
NS500138:372:H7K35BGXC:1:12111:21030:5973	16	chr2L	20634	31	36M	*	0	0		CTGGTCGGTACTAATATTTTTCTAAGACTGTATAC	EEEE/EEEEEEEEEEEEEEEEEEEEEEEEAAAAA
NS500138:372:H7K35BGXC:1:11205:21596:8754	0	chr2L	20690	1	36M	*	0	0		CGTAGGTATGTATCCGTGAGTAAGCTGAGGTGCAAA	AAAAEEEEEEEEEEEEEEEEEEEEEEEEEEEE
NS500138:372:H7K35BGXC:1:11201:23277:9037	0	chr2L	20726	1	36M	*	0	0		ACCGTATCACCATTGCTTTTCGGCTGATAACCAAT	AAAAEEEEEEEEEEEEEEEEEEEEEEEEEEEE
NS500138:372:H7K35BGXC:1:11211:25328:12582	0	chr2L	20760	1	33M	*	0	0		ATACCAAGTAAGTCAACTGTTTCCCAAATGTAT	AAAAEEEEEEEEEEEEEEEEEEEEEEEEEEEE
						rnext	pnext	tlen			

- Bowtie2 ‘map quality’: a complex subject.
 - 0 and 1 indicate that a read maps to multiple locations in the genome and the position cannot be uniquely assigned. Whether 0 or 1 depends on how many mismatches there are to the genome assembly.
 - A score of 42 is a perfect unique match to the genome assembly
 - Values in between are probably mapped to the right location, but there are other potential matches. Here, the score is roughly $-10 * \log_{10}(\text{probability of mismatch})$
- For a great discussion about Bowtie2 map quality scoring, please see:
- <http://biofinysics.blogspot.com/2014/05/how-does-bowtie2-assign-mapq-scores.html>

The .sam/.bam format

Example: single end mapping to *Drosophila melanogaster*

“qname” = Name of the read from fastQ	flag	rname	pos	mapq	cigar					Sequence	Phred Quality Scores
NS500138:372:H7K35BGXC:1:12111:21030:5973	16	chr2L	20634	31	36M	*	0	0		CTGGTCGGTACTAATATTTTTCTAAGACTGTATAC	EEEE/EEEEEEEEEEEEEEEEEEEEEEEEAAAAA
NS500138:372:H7K35BGXC:1:11205:21596:8754	0	chr2L	20690	1	36M	*	0	0		CGTAGGTATGTATCCGTGAGTAAGCTGAGGTGCAAA	AAAAEEEEEEEEEEEEEEEEEEEEEEEEEEEEEE
NS500138:372:H7K35BGXC:1:11201:23277:9037	0	chr2L	20726	1	36M	*	0	0		ACCGTATCACCATTGCTTTTCGGCTGATAACCAAT	AAAAEEEEEEEEEEEEEEEEEEEEEEEEEEEEEE
NS500138:372:H7K35BGXC:1:11211:25328:12582	0	chr2L	20760	1	33M	*	0	0		ATACCAAGTAAGTCAACTGTTCCCAAATGTAT	AAAAEEEEEEEEEEEEEEEEEEEEEEEEEEEEEE

rnext pnext tlen

- Flags are equally confusing at first. Most are designated for describing paired-end data.
- In the example above, we have flags of 16 and 0
- 16 means “minus strand”, 0 means “plus strand” mapping.
- <https://broadinstitute.github.io/picard/explain-flags.html>

Binary (Decimal)	Hex	Description
00000000001 (1)	0x1	Is the read paired?
00000000010 (2)	0x2	Are both reads in a pair mapped “properly” (i.e., in the correct orientation with respect to one another)?
00000000100 (4)	0x4	Is the read itself unmapped?
00000001000 (8)	0x8	Is the mate read unmapped?
00000010000 (16)	0x10	Has the read been mapped to the reverse strand?
00000100000 (32)	0x20	Has the mate read been mapped to the reverse strand?
00001000000 (64)	0x40	Is the read the first read in a pair?
00010000000 (128)	0x80	Is the read the second read in a pair?
00100000000 (256)	0x100	Is the alignment not primary? (A read with split matches may have multiple primary alignment records.)
01000000000 (512)	0x200	Does the read fail platform/vendor quality checks?
10000000000 (1024)	0x400	Is the read a PCR or optical duplicate?

The .sam/.bam format

Example: single end mapping to *Drosophila melanogaster*

“qname” = Name of the read from fastQ	flag	rname	pos	mapq	cigar					Sequence	Phred Quality Scores
NS500138:372:H7K35BGXC:1:11111:25532:10062	16	chr2L	356378	42	35M	*	0	0		CTGTGGCCGTACTCTGACTGCATAAGCACTGTAAC	EEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEAAAA
NS500138:372:H7K35BGXC:1:11203:22330:18504	1040	chr2L	356378	42	35M	*	0	0		CTGTGGCCGTACTCTGACTGCATAAGCACTGTAAC	EEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEAAAA
NS500138:372:H7K35BGXC:1:11105:14496:15080	1024	chr2L	356397	42	36M	*	0	0		GCATAAGCACTGTAAGTCCGGATAGACAGACAGGG	AA6AA//AEE/EEAE/E//EA//EEE/AAEEEEEE/E
NS500138:372:H7K35BGXC:1:11109:5988:6069	0	chr2L	356397	42	36M	*	0	0		GCATAAGCACTGTAAGTCCGGATAGACAGACAGGG	AAAAEEEEEEEEEEEEEEEEEEEE6EEEEEEEEEE

rnext pnext tlen

- Flags are equally confusing at first. Most are designated for describing paired-end data.
- Flag values **sum**, the above example shows two duplicated reads, one on the minus, the other on the plus strand.
- <https://broadinstitute.github.io/picard/explain-flags.html>

Binary (Decimal)	Hex	Description
00000000001 (1)	0x1	Is the read paired?
00000000010 (2)	0x2	Are both reads in a pair mapped “properly” (i.e., in the correct orientation with respect to one another)?
00000000100 (4)	0x4	Is the read itself unmapped?
00000001000 (8)	0x8	Is the mate read unmapped?
00000010000 (16)	0x10	Has the read been mapped to the reverse strand?
00000100000 (32)	0x20	Has the mate read been mapped to the reverse strand?
00001000000 (64)	0x40	Is the read the first read in a pair?
00010000000 (128)	0x80	Is the read the second read in a pair?
00100000000 (256)	0x100	Is the alignment not primary? (A read with split matches may have multiple primary alignment records.)
01000000000 (512)	0x200	Does the read fail platform/vendor quality checks?
10000000000 (1024)	0x400	Is the read a PCR or optical duplicate?

The .sam/.bam format

Example: paired-end mapping to *Drosophila melanogaster*

“qname” = Name of the read from fastQ	flag	rname	pos	mapq	cigar					Sequence	Phred Quality Scores
NS500138:372:H7K35BGXC:1:12201:14062:6164	147	chr2L	83234	42	36M	=	83017	-253		CTAAAAAAGTAAGTTATCCACGAAACTGGACAGATT	EEEEEEEEAAEEEE///AEEEEEEEEEE6EEEEAAAAA
NS500138:372:H7K35BGXC:1:11106:22617:13759	147	chr2L	83253	42	36M	=	83155	-134		ACGAAACTGGACAGATTACTCGTCTTTCTTTTCGTTT	EEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEAAAAA
NS500138:372:H7K35BGXC:1:11105:4862:1314	163	chr2L	83398	42	36M	=	83407	43		AGTATATAGGAACACGCTACCCGAAGAACGGAGAGT	AAAAAEEEEEEAEEEEEEEEEEEEEEEE/EEEEEEEE
NS500138:372:H7K35BGXC:1:11105:4862:1314	83	chr2L	83407	42	34M	=	83398	-43		GAACACGCTACCCGAAGAACGGAGAGTTTGAAA	EEEEEEAEEEEEEEEEEEEEEEEEEEE6EEAAAAA

rnext pnext tlen

- This paired-end example fills in more columns. The flags are also more complex.
- 147 = read paired, properly, reverse strand, second read.
- 163 = read paired, properly, plus strand, second read.
- 83 = read paired, properly, minus strand, first read

Binary (Decimal)	Hex	Description
00000000001 (1)	0x1	Is the read paired?
00000000010 (2)	0x2	Are both reads in a pair mapped “properly” (i.e., in the correct orientation with respect to one another)?
00000000100 (4)	0x4	Is the read itself unmapped?
00000001000 (8)	0x8	Is the mate read unmapped?
00000010000 (16)	0x10	Has the read been mapped to the reverse strand?
00000100000 (32)	0x20	Has the mate read been mapped to the reverse strand?
00001000000 (64)	0x40	Is the read the first read in a pair?
00010000000 (128)	0x80	Is the read the second read in a pair?
00100000000 (256)	0x100	Is the alignment not primary? (A read with split matches may have multiple primary alignment records.)
01000000000 (512)	0x200	Does the read fail platform/vendor quality checks?
10000000000 (1024)	0x400	Is the read a PCR or optical duplicate?

Next week, we will learn how to import these data to R

These features (flag, mapq) can be used to filter out bad stuff.

- First, we need to learn how to generate .bam files from .fastq data.
- Lots of work to set up the first time... trial and error... once you have the data, you won't think about this part very often.
- But you need to be systematic and the mapping needs to be reproducible.

Fast-forward to the day you submit your paper

Public Database Submission

- Every published genomics dataset must be deposited in a public repository (e.g., Gene Expression Omnibus/Sequence Read Archive).
- What does GEO/SRA want?
 - Original FastQ files (trimmed reads might be OK), gzipped.
 - All pertinent information about how you processed the files
 - None of the .bam files.

Fast-forward to the day you submit your paper

Materials and Methods

- We all know there are good and bad Materials and Methods sections in papers. **You should all commit to writing really good ones.**
- Your analysis should be perfectly reproducible. Because GEO doesn't take .bam files, and this is the centerpiece of an analysis, you also need to document exactly how you got your reads mapped.
- It is also good (and sometimes necessary) to enumerate how many reads you started with, what the mapping rate was, what the duplication rate was, and so on...

Reproducible genomics begins with reproducible code

- I have made two ‘test scripts’ that will map a test dataset *if* the computing environment is set up properly.
- It also highlights one possible way to set up a single script that you could use for all your mapping needs for either paired- or single-end data

Reproducible approaches require reproducible directory structures.

- My suggested directory structure for one sequencing run is:

>base_directory (e.g., 211018_pilot_experiment)

>Raw_Data

>Trimmed_Reads

>Mapped_Reads

- outlog

- errlog

- mapping_script.sh *(not the master version, but the version that you use for this particular mapping)*

The Test Script

- I have created test scripts for both paired- and single-end data. They are in the “Quest” directory in the class Github Repo.
- I have created an annotated markdown that works through the paired-end script line by line.
- I have also included documentation of how I set up the class Quest allocation to run this code.

```
1  #!/bin/bash
2  #SBATCH --account b1042
3  #SBATCH --partition genomicsquestA
4  #SBATCH --nodes=1
5  #SBATCH --ntasks-per-node=12
6  #SBATCH --time=00:20:00
7  #SBATCH --mem-per-cpu=12G
8  #SBATCH --job-name=test_script_PE          ## change this to job name
9  #SBATCH --output=outlog                    ## will write files to the current dir.
10 #SBATCH --error=errlog
11
12 basedir=</full/path/to/the/project/directory>    ## change to full path
13
14 ## make destinations
15 mkdir ${basedir}/Trimmed_Reads
16 mkdir ${basedir}/Mapped_Reads
17
18 ## load modules be verbose about versions since the defaults can change!
19 module load python/3.8.4
20 module load fastqc/0.11.5
21 module load bowtie2/2.4.1
22 module load samtools/1.10.1
23 module load picard/2.21.4
24
25 ## assign the variables for the filepaths for this mapping
26 rawdir=Raw_Data
27 trimdir=Trimmed_Reads
28 mapdir=Mapped_Reads
29 index=/projects/b1059/RforBiologists/Bowtie_Indices/dm6/dm6 ## make sure you have correct index here
30
31 #####
```


Goals/Plans

- **Today, Monday October 18:** Set up directory structures for running the test scripts. Upload data to your Quest folder. Modify the test script(s) to run correctly. See if you can get them to run.
- **By next Monday:** Upload your own raw fastQ data to project directories you create in your Quest folder. Modify the test script appropriately, and map your data to the appropriate reference genome.
- **Next class (Thursday) will be entirely project-time, figuring out how to do the mapping. I will be available for consultations.**

Big hints:

- Please read the script annotation document as well as the comments I've added to the script itself.
- When you initially upload the script to Quest, **it will not be executable**. This is because the correct permission has not been set. This is a frequent thing to encounter. **You fix this by running the following bash command:**

```
$  
$ chmod -x my_script_name.sh
```

- (On the assumption that you are *in the directory that contains the script* and that the name of the script is “my_script_name.sh”).

Take the time now to find the scripts and download them to your computer.

Open them in your code-aware text editor.

I will then give you an overview of the workflow.

General Workflow

